

# COMP0188

# Deep Representation and Learning

Joshua Spear  
[joshua.spear.21@ucl.ac.uk](mailto:joshua.spear.21@ucl.ac.uk)

# Today

- Reviewing VAEs
- Brief introduction to tokenisation in NLP
- Attention
- Transformers
- Diffusion models

# VAE update

# VI with parameter learning

- ELBO:

$$\text{ELBO}(q_\phi(z|x)) = -KL(q_\phi(z|x)||p_\theta(z)) + \mathbb{E}_{q_\phi(z|x)} [\log p_\theta(x|z)]$$

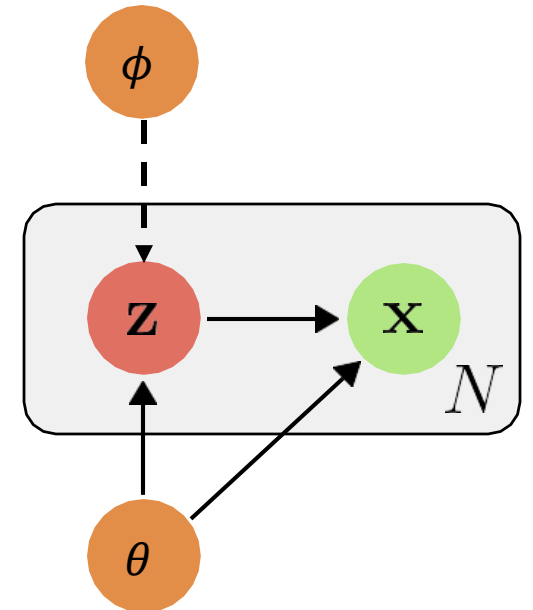
- Optimise jointly with respect to the model parameters  $\theta$  and the variational parameters  $\phi$

- Intuitively:

- $KL(q_\phi(z|x)||p_\theta(z))$  is the divergence between the variational **posterior** and the **prior** over the latent variables and;
- $\mathbb{E}_{q_\phi(z|x)}[\log p_\theta(x|z)]$  is the expected log likelihood of the **reconstruction**

- Therefore, VAEs are **regularized autoencoders** where the form of the regulariser is **defined by the prior**!

- $p_\theta(z)$  is fixed and defined as  $p_\theta(z) \sim N(0,1)$  (for this discussion)



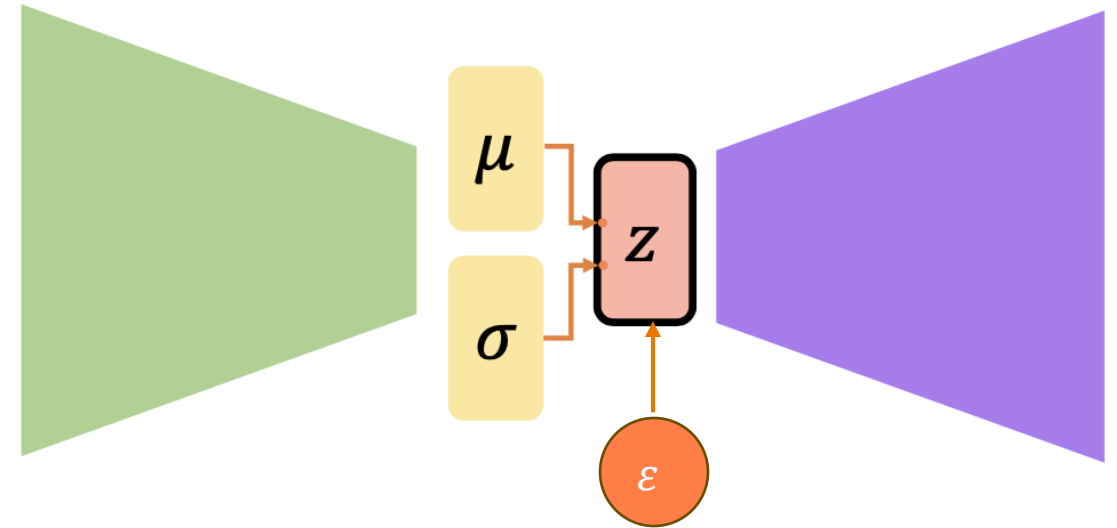
# Reparameterisation trick

- Recall

$$p_{\theta}(z|x) = N(\mu, \sigma^2)$$

- Instead assume

$$p_{\theta}(z|x) = \mu + \sigma^2 \odot N(0,1)$$

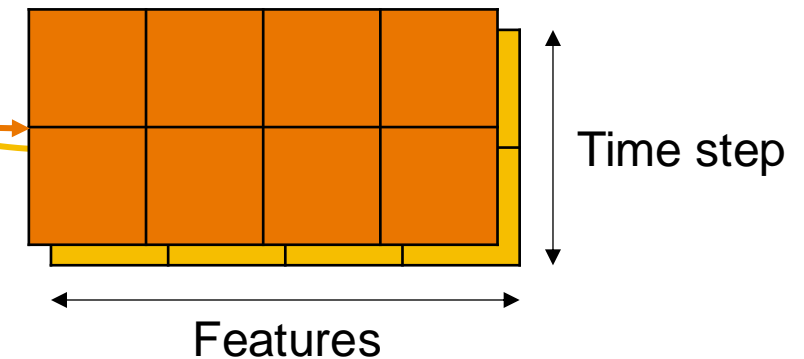


# Sequence modelling with RNNs: A review

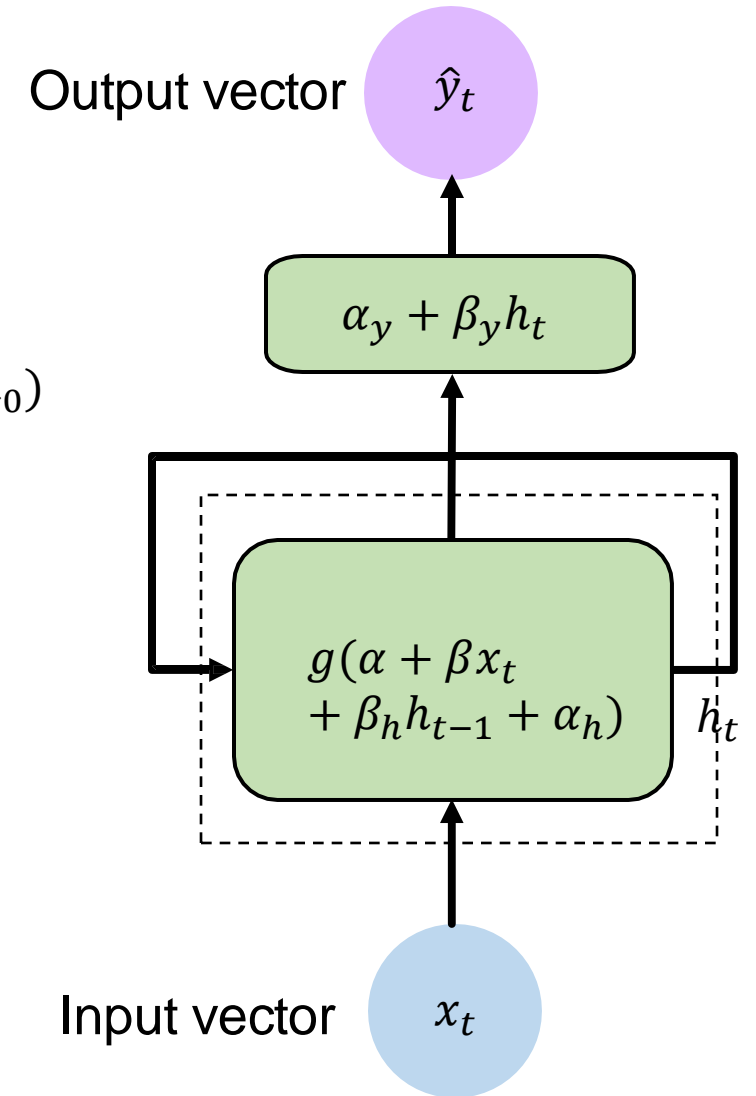
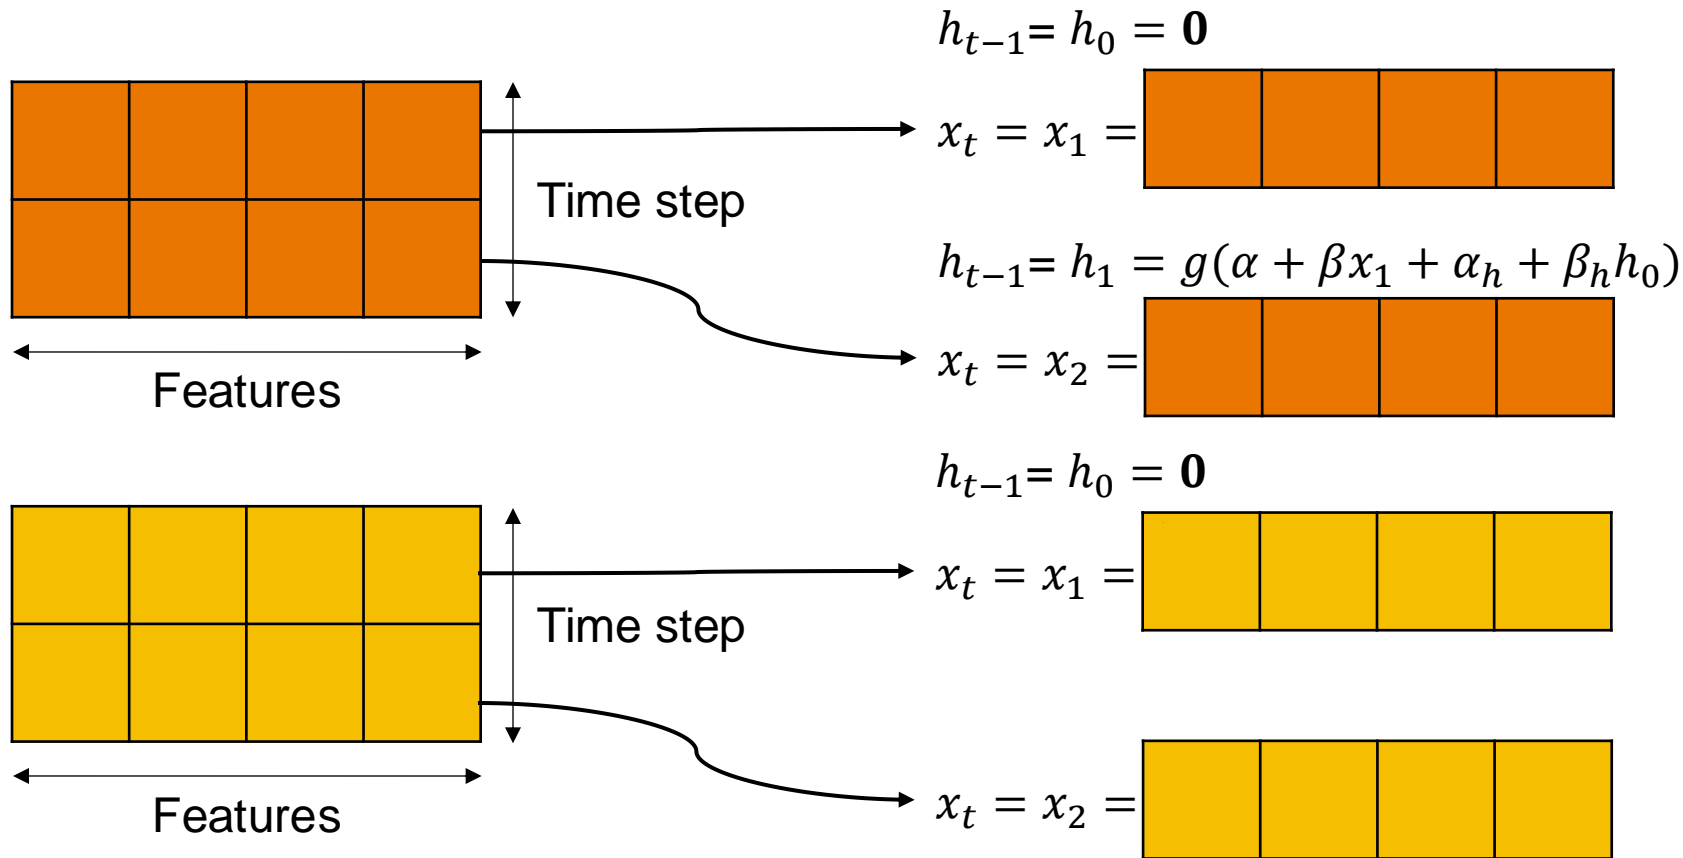
# Timeseries tensor shapes

- Similarly to (black and white) images, timeseries add another dimension to the input
- For time series the input dimension will be:
  - (batch size, time step, feature)

	date	meantemp	humidity	wind_speed	meanpressure
0	2013-01-01	10.000000	84.500000	0.000000	1015.666667
1	2013-01-02	7.400000	92.000000	2.980000	1017.800000
2	2013-01-03	7.166667	87.000000	4.633333	1018.666667
3	2013-01-04	8.666667	71.333333	1.233333	1017.166667
4	2013-01-05	6.000000	86.833333	3.700000	1016.500000
5	2013-01-06	7.000000	82.800000	1.480000	1018.000000
6	2013-01-07	7.000000	78.600000	6.300000	1020.000000
7	2013-01-08	8.857143	63.714286	7.142857	1018.714286
8	2013-01-09	14.000000	51.250000	12.500000	1017.000000
9	2013-01-10	11.000000	62.000000	7.400000	1015.666667



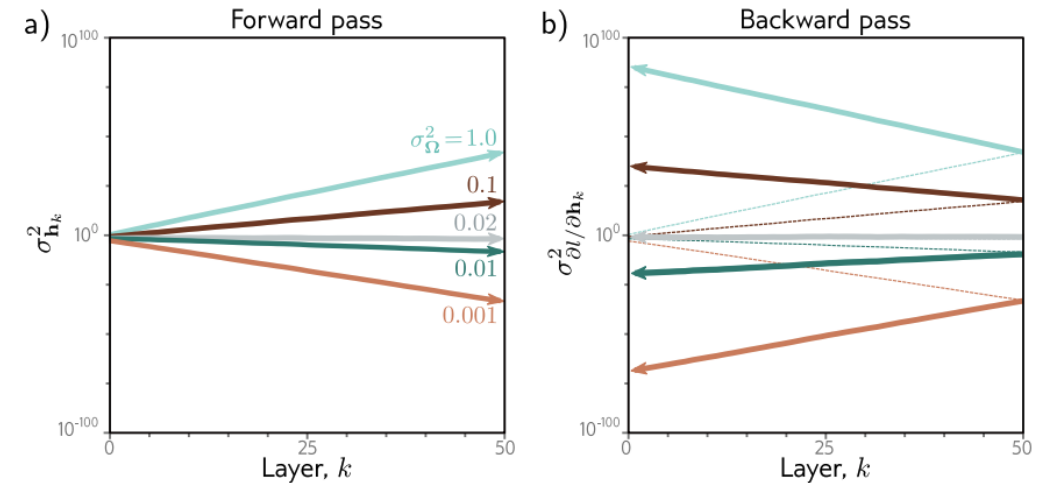
# RNN





# Vanishing/exploding gradients

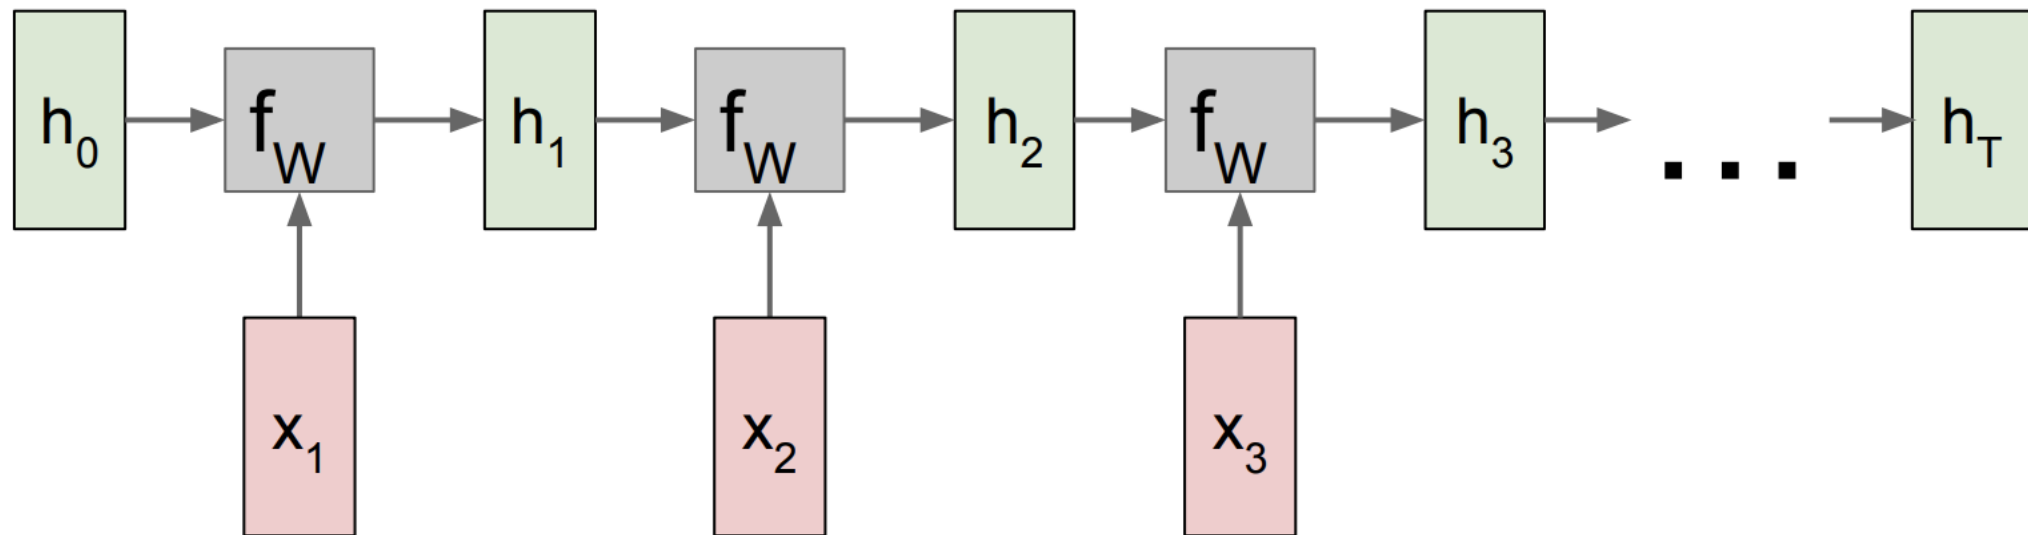
- Vanishing and exploding gradients occur when weight updates become **too small** or **too large**
- RNNs can be viewed as (broadly speaking), a very deep MLP where:
  - The time horizon denotes the “depth”
  - The same weight matrix,  $W$ , is used at each layer
- Since the **same** weight matrix is used, initialisation strategies such as He and Xavier that rely on independent weights don't apply!



We saw an example of this in **week 2** where, if the **weights were inappropriately initialised**, the **variance of the hidden state** could **explode/vanish**

# Vanishing/exploding gradients

- Vanishing gradients are **desirable to some extent** as it is reasonable to assume that for timestep  $t$ , information at timestep  $k: k < t$  is *more useful* than information at timestep  $k': k'' < k$ 
  - Is there a more **selective** way to learn historical relationships... (hint: yes! **Transformers**)?



# Attention is all you need...

- Does not require bootstrapping of predictions
  - Timesteps can be processed in parallel
- Handles long term dependencies more effectively

---

## Attention Is All You Need

---

**Ashish Vaswani\***  
Google Brain  
avaswani@google.com

**Noam Shazeer\***  
Google Brain  
noam@google.com

**Niki Parmar\***  
Google Research  
nikip@google.com

**Jakob Uszkoreit\***  
Google Research  
usz@google.com

**Llion Jones\***  
Google Research  
llion@google.com

**Aidan N. Gomez\* †**  
University of Toronto  
aidan@cs.toronto.edu

**Łukasz Kaiser\***  
Google Brain  
lukaszkaizer@google.com

**Illia Polosukhin\* †**  
illia.polosukhin@gmail.com

# Language embeddings (A brief introduction to tokenisation)

# Numerical representations

- Machine learning models require **numerical representations**
- Defining **good** numerical representations is the ultimate goal of representation learning for language

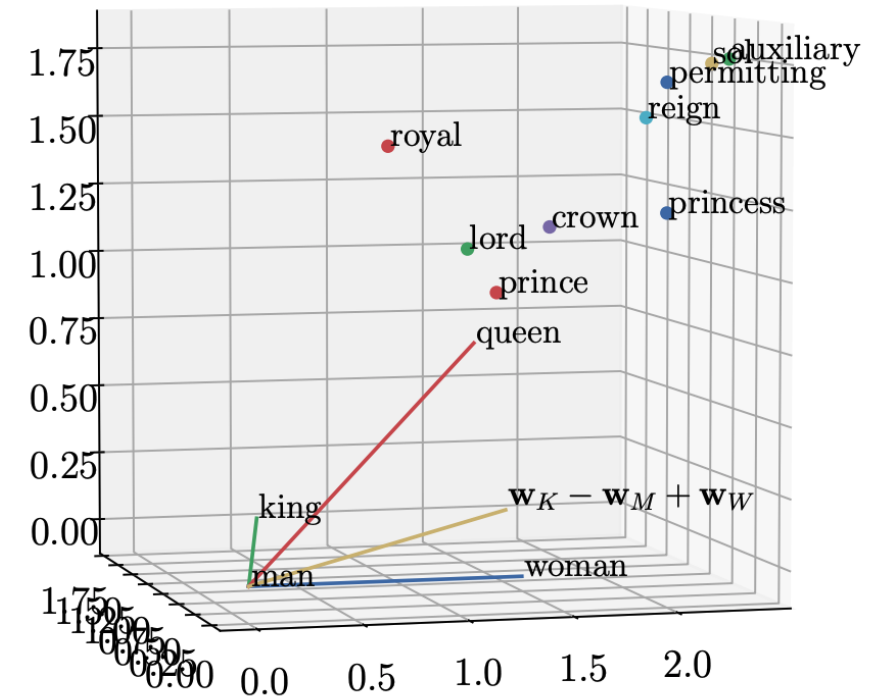


Figure 1. The relative locations of word embeddings for the analogy "man is to king as woman is to ..?". The closest embedding to the linear combination  $\mathbf{w}_K - \mathbf{w}_M + \mathbf{w}_W$  is that of *queen*. We explain why this occurs and interpret the difference between them.

# Tokenisation

- What defines a single input?
  - “The quick brown...”, [“The”, “quick”, “brown”, ...], [“T”, “h”, “e”, “q”, “u”, “i”, “c”, “k”, ...]
- Defining a **single input** is known as **tokenization**
  - Broadly speaking, optimal tokenisation defines **semantically meaningful** substrings of text
- BERT tokenizer (Word Piece)
  - “The quick brown fox” → [“The”, “quick”, “brown”, “fox”]
  - “Why the edits made under my username Hardcore Metallica Fan were reverted?” → [“Why”, “the”, “edit”, “##s”, “made”, “under”, “my”, “user”, “##name”, “Hard”, “##core”, “Metal”, “##lica”, “Fan”, “were”, “reverted”, “?”]
- [https://www.youtube.com/watch?v=9vM4p9NN0Ts&t=2167s&ab\\_channel=StanfordOnline](https://www.youtube.com/watch?v=9vM4p9NN0Ts&t=2167s&ab_channel=StanfordOnline)

# Token embeddings

- Each embedding is mapped to a numeric ID:
  - “The quick brown fox” → {“The”: 1109, “quick”: 3613, “brown”: 3058, “fox”: 17594}
- Each numeric ID indexes a vector
  - torch.nn.Embedding is a trainable weight matrix

```
class TokenEmbedding(nn.Module):  
  
    def __init__(  
        self,  
        vocab_size:int,  
        embed_size:int,  
    ):  
        super().__init__()  
        self.embed_size = embed_size  
        # (m, seq_len) --> (m, seq_len, embed_size)  
        # padding_idx is not updated during training, remains as fixed pad (0)  
        self.token = torch.nn.Embedding(vocab_size, embed_size, padding_idx=0)  
  
    def forward(  
        self,  
        sequence:Float[tuple(torch.Tensor, "batch_size max_length")]  
    ) -> Float[tuple(torch.Tensor, "batch_size max_length d_model")]:  
        _token_embed = self.token(sequence)  
        x = _token_embed  
        return x  
  
_tmp_te = TokenEmbedding(  
    vocab_size=len(tokenizer.vocab),  
    embed_size=32  
)  
for i in _tmp_te.token.named_parameters():  
    print(i[0])  
    print(i[1].shape)
```

✓ 0.0s

# Attention



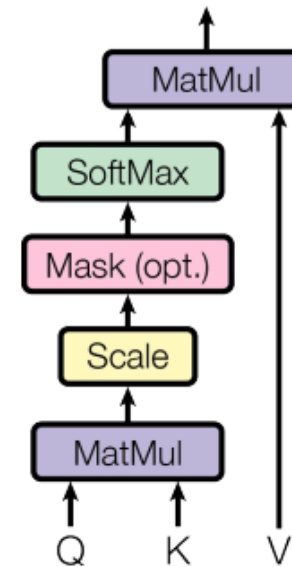
# Attention intuition



# Scaled dot product attention

- Proposed in “Attention is all you need”
- Attention intuition:
  - Provide a set of input queries,  $Q$
  - Map these to a set of (input) keys,  $K$ 
    - Broadly speaking by computing how “close” the query is to a key, we are computing how “relevant” the key is to the query
  - Given some values  $V$  with one-to-one mapping to keys, up/down weight the magnitude of  $V$  by the previously calculated relevance

Scaled Dot-Product Attention



# Scaled dot product attention

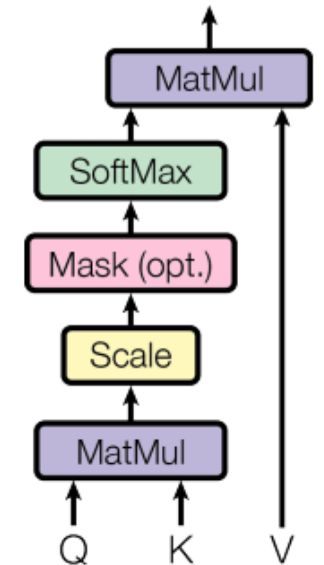
- Process:

1.  $\text{out}_1 = QK^T$ 
  - $Q \in \mathbb{R}^{d_Q}, K \in \mathbb{R}^{d_K} \Rightarrow \text{out}_1 \in \mathbb{R}$
2.  $\text{out}_2 = \frac{\text{out}_1}{\sqrt{d_K}} \in \mathbb{R}$
3.  $\text{out}_3 = \text{softmax}(\text{out}_2) \in \mathbb{R}$
4.  $\text{out}_4 = \text{out}_3 V$ 
  - $V \in \mathbb{R}^{d_V} \Rightarrow \text{out}_4 \in \mathbb{R}^{d_V}$

- In words:

- Given a single query vector  $Q$ , and **single** key value pair  $K, V$
- How much attention should be paid to the value vector of the associated key

Scaled Dot-Product Attention



# Scaled dot product attention

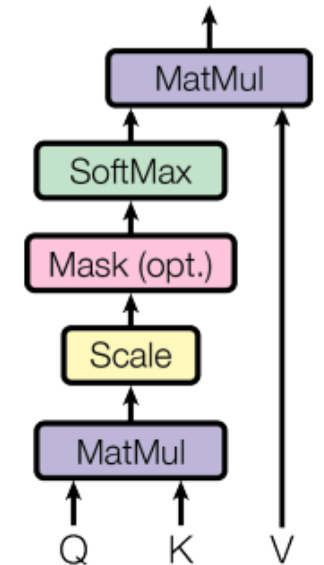
- Process:

- $\text{out}_1 = QK^T$  ← Dot product can be considered a **measure of similarity**
  - $Q \in \mathbb{R}^{d_Q}, K \in \mathbb{R}^{d_K} \Rightarrow \text{out}_1 \in \mathbb{R}$
- $\text{out}_2 = \frac{\text{out}_1}{\sqrt{d_K}} \in \mathbb{R}$
- $\text{out}_3 = \text{softmax}(\text{out}_2) \in \mathbb{R}$  ← Given the similarity (defined in the softmax), how much should the **associated value be weighted**
- $\text{out}_4 = \text{out}_3 V$ 
  - $V \in \mathbb{R}^{d_V} \Rightarrow \text{out}_4 \in \mathbb{R}^{d_V}$

- In words:

- Given a single query vector  $Q$ , and **single** key value pair  $K, V$
- How much attention should be paid to the value vector of the associated key
- A single query and key-value is not so helpful...

Scaled Dot-Product Attention



# Scaled dot product attention

- Process for multiple queries and keys:

- $out_1 = QK^T$

- $Q \in \mathbb{R}^{n_Q \times d_K}, K \in \mathbb{R}^{n_K \times d_K} \Rightarrow out_1 \in \mathbb{R}^{n_Q \times n_K}$

- $out_2 = \frac{out_1}{\sqrt{d_K}} \in \mathbb{R}^{n_Q \times n_K}$

- $out_3 = \text{softmax}(out_2) \in \mathbb{R}^{n_Q \times n_K}$

- $out_4 = out_3 V$

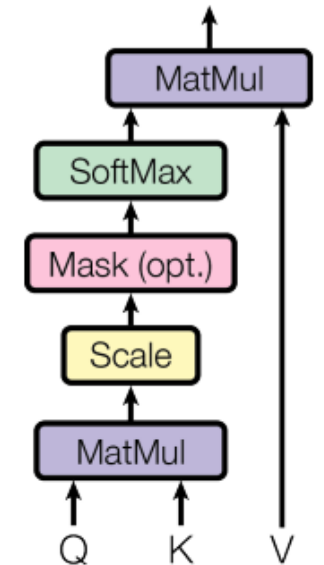
- $V \in \mathbb{R}^{n_K \times d_V} \Rightarrow out_4 \in \mathbb{R}^{n_Q \times d_V}$

For each query ( $n_Q$  of them) there now exists a dot product for each key ( $n_K$  of them)

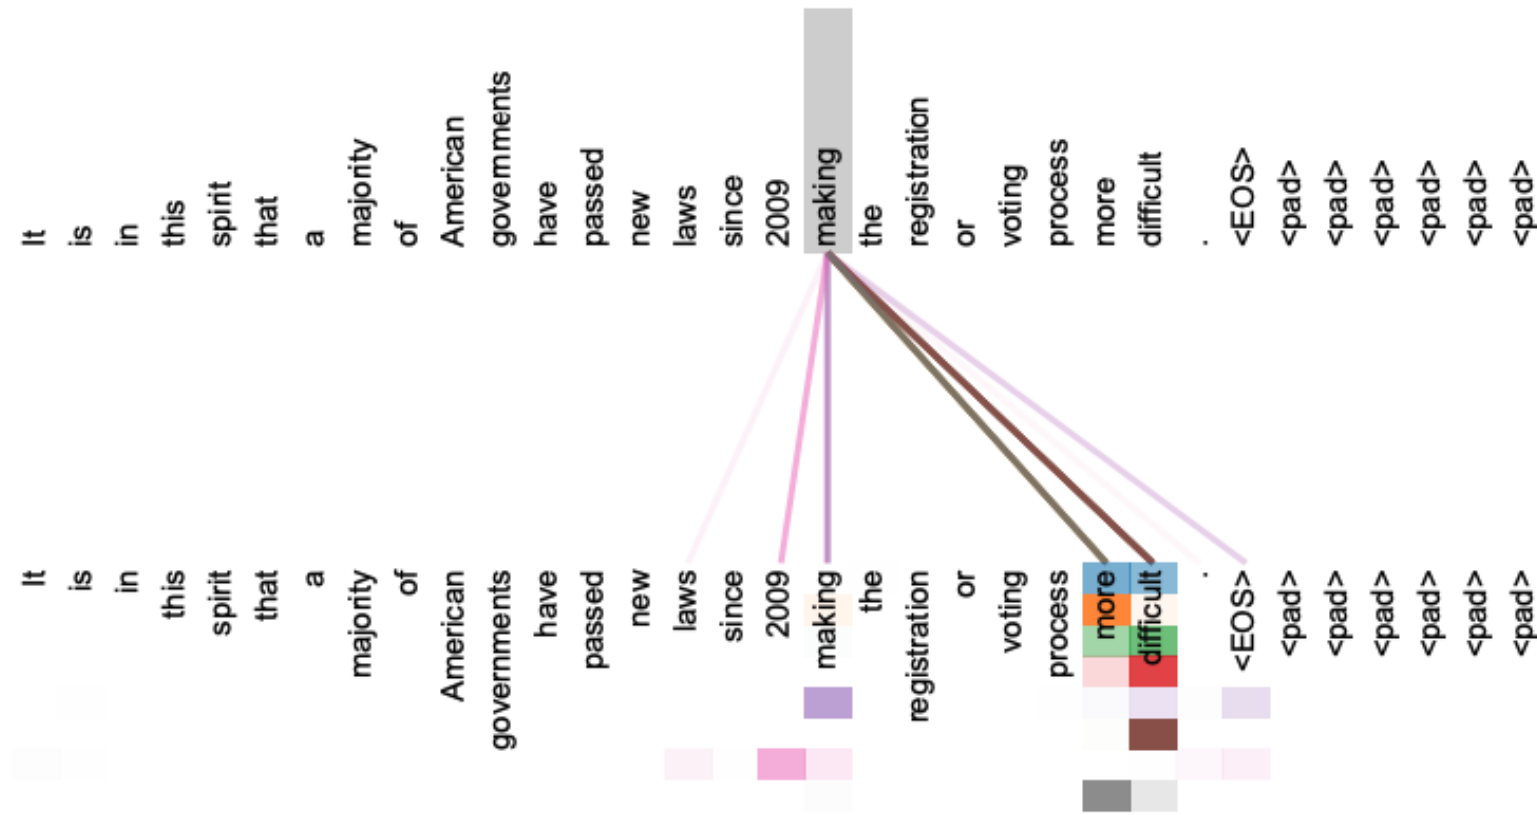
- Now, for each query ( $n_Q$  of them) there is a  $d_V$  associated “answer” vector

- $V \in \mathbb{R}^{n_K \times d_V} \Rightarrow$  The same number of keys and values are required

Scaled Dot-Product Attention



# Self attention



# Self attention

- Scaled dot product attention:  $\text{out}_4 = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$
- In **self** attention,  $Q, K$  and  $V$  are **all** derived from the same input i.e., **self**
- Given an input, ["The", "quick", ...]
  - Tokenize and convert to embeddings via lookup
  - $X = \begin{matrix} 0.5 & 0.2 \\ -1.2 & -1.4 \end{matrix}$  where  $X \in \mathbb{R}^{d_s \times d_x}$  and  $x$  is a row vector from  $X$
  - $Q = x\ddot{W}^Q, K = x\ddot{W}^K, V = x\ddot{W}^V$  such that  $W^Q \in \mathbb{R}^{d_x \times d_K}, W^K \in \mathbb{R}^{d_x \times d_K}, W^V \in \mathbb{R}^{d_x \times d_V}$
- $W^Q, W^K$  and  $W^V$  are **learnable** parameter matrices

# Self attention

- Given an input, [“The”, “quick”, ...]
  - Tokenize and convert to embeddings via lookup
 

$$\begin{matrix} 0.5 & 0.2 \\ -1.2 & -1.4 \\ \dots & \dots \end{matrix}$$
  - $X = \begin{matrix} 0.5 & 0.2 \\ -1.2 & -1.4 \\ \dots & \dots \end{matrix}$  where  $X \in \mathbb{R}^{d_S \times d_X}$  and  $x$  is a row vector from  $X$ 
    - $d_S$  defines the “sequence length” i.e., the number of tokens in the input
    - $d_X$  defines the vector dimension for each token
  - Generally, set  $d_Q = d_K = d_V$
  - $q = xW^Q$ ,  $k = xW^K$ ,  $v = xW^V$  such that  $W^Q \in \mathbb{R}^{d_X \times d_Q}$ ,  $W^K \in \mathbb{R}^{d_X \times d_K}$ ,  $W^V \in \mathbb{R}^{d_X \times d_V}$ 
    - **IMPORTANT:** Linear projections are applied **individually** to each element of  $X$
    - $q \in \mathbb{R}^{1 \times d_Q} \Rightarrow Q \in \mathbb{R}^{d_S \times d_Q}$  etc.

- Recall:

$$\text{out}_4 = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

- $QK^T \in \mathbb{R}^{d_S \times d_S} \Rightarrow$ 
  - $QK^T$  defines the **relevance** between **every combination** of input token:

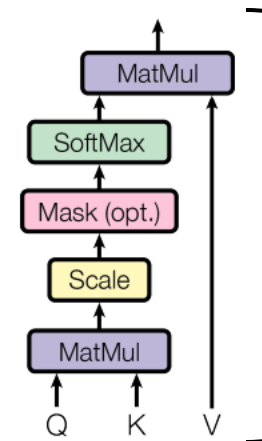
The × quick



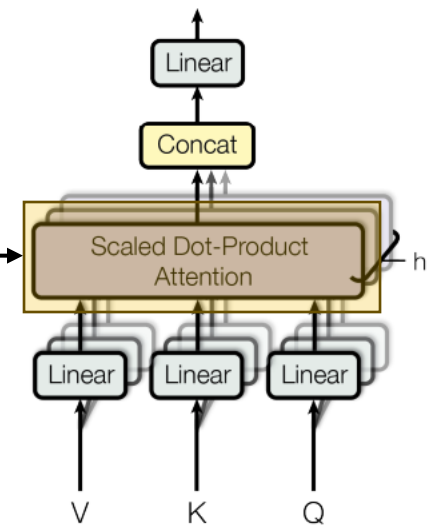
# Multi-head attention

- $h$  attention heads are computed in parallel
- Process
  - The dimensions of  $Q, K$  and  $V$  are **split** across the  $h$  heads
  - Each head outputs a matrix of dimension  $\mathbb{R}^{d_s \times d_v/h}$
  - This means for a given query-key pair, **different softmax values** (amounts of attention) can be applied to different dimensions of  $V$
- $\text{out} = \text{concat}(\text{head}_1, \text{head}_2, \dots) W^O$  where:
  - $\text{head}_i = \text{softmax}\left(\frac{Q_i K_i}{\sqrt{d_k}}\right) V_i, W^O \in \mathbb{R}^{h d_v \times d_o} \Rightarrow \text{out} \in \mathbb{R}^{d_s \times d_o}$
  - $h$  is the number of scaled dot-product attention cells
  - $d_o = d_x$

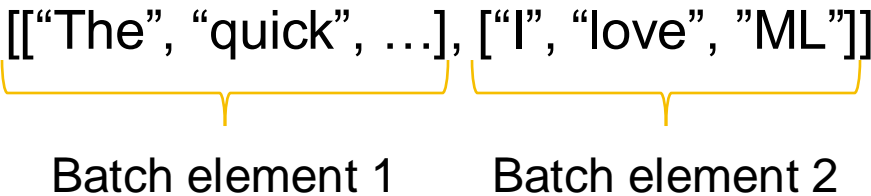
Scaled Dot-Product Attention



Multi-Head Attention



# Batched self attention

- Since this is SGD, we perform computations over batches
- Within batch computations are **independent**  $\Rightarrow$ 
  - The previous (multi head) self attention process is applied to each element of the batch **independently**
- Batch size 2:
  - 
    - Batch element 1
    - Batch element 2

# Multi-head self attention (implementation)

```
class MultiHeadSelfAttention(nn.Module):

    def __init__(
        self,
        d_model:int,
        n_heads:int = 1
    ) -> None:
        super().__init__()
        assert d_model % n_heads == 0
        self.__n_heads = n_heads
        self.__d_v = self.__d_k = int(d_model/n_heads)
        self.W_q = nn.Linear(d_model, d_model, bias=False)
        self.W_k = nn.Linear(d_model, d_model, bias=False)
        self.W_v = nn.Linear(d_model, d_model, bias=False)
        self.__norm = torch.sqrt(torch.tensor(d_model))
        self.soft = nn.Softmax()
        self.W_o = nn.Linear(d_model, d_model, bias=False)
```


```
def forward(
    self,
    x:Float[torch.Tensor, "batch_size max_length d_model"],
):
    # Linear regression is applied to only the last dimension i.e.,
    # this is equivalent to independantly applying the regression to each
    # element of each batch independently
    # Output dimension:
    # Float[torch.Tensor, "batch_size max_length d_model"]
    batch_size, seq_len, embed_dim = x.shape
    q = self.W_q(x)
    k = self.W_k(x)
    v = self.W_v(x)
    # Split up by heads ->
    #(batch_size, max_length, heads, embedding_dim/heads)
    # permute ready for matmul ->
    #(batch_size, heads, max_length, embedding_dim/heads)
    q = q.view(
        batch_size, seq_len, self.__n_heads, self.__d_k
    ).permute(0,2,1,3)
    k = k.view(
        batch_size, seq_len, self.__n_heads, self.__d_k
    ).permute(0,2,1,3)
    v = v.view(
        batch_size, seq_len, self.__n_heads, self.__d_v
    ).permute(0,2,1,3)
    _t_1 = torch.matmul(q, torch.transpose(k, dim0=2, dim1=3))
    _t_2 = self.soft(_t_1/self.__norm)
    _t_3 = torch.matmul(_t_2, v)
    _t_4 = _t_3.permute(0,2,1,3).reshape(
        batch_size, seq_len, embed_dim
    )
    return self.W_o(_t_4)
```

# Multi-head attention and long-term dependencies

- RNNs suffer:
  - From **vanishing gradients**  $\Rightarrow$  struggle to model **long term dependencies**
  - **Slow computation** due to the **autogressive** processing of input data
    - Sequential operations scale as  $O(n)$  where  $n$  is the sequence length
- Multi-head attention addresses both:
  - **Long term dependencies**: Modelled via attention. Don't need to go through  $x_{t-1}, \dots, x_{t-k}$  to model  $g(x_t, x_{t-k-1})$
  - **Slow computation**: The dependencies between  $x_t, \dots, x_1$  are modelled **simultaneously**
    - Sequential operations scale as  $O(1)$ !

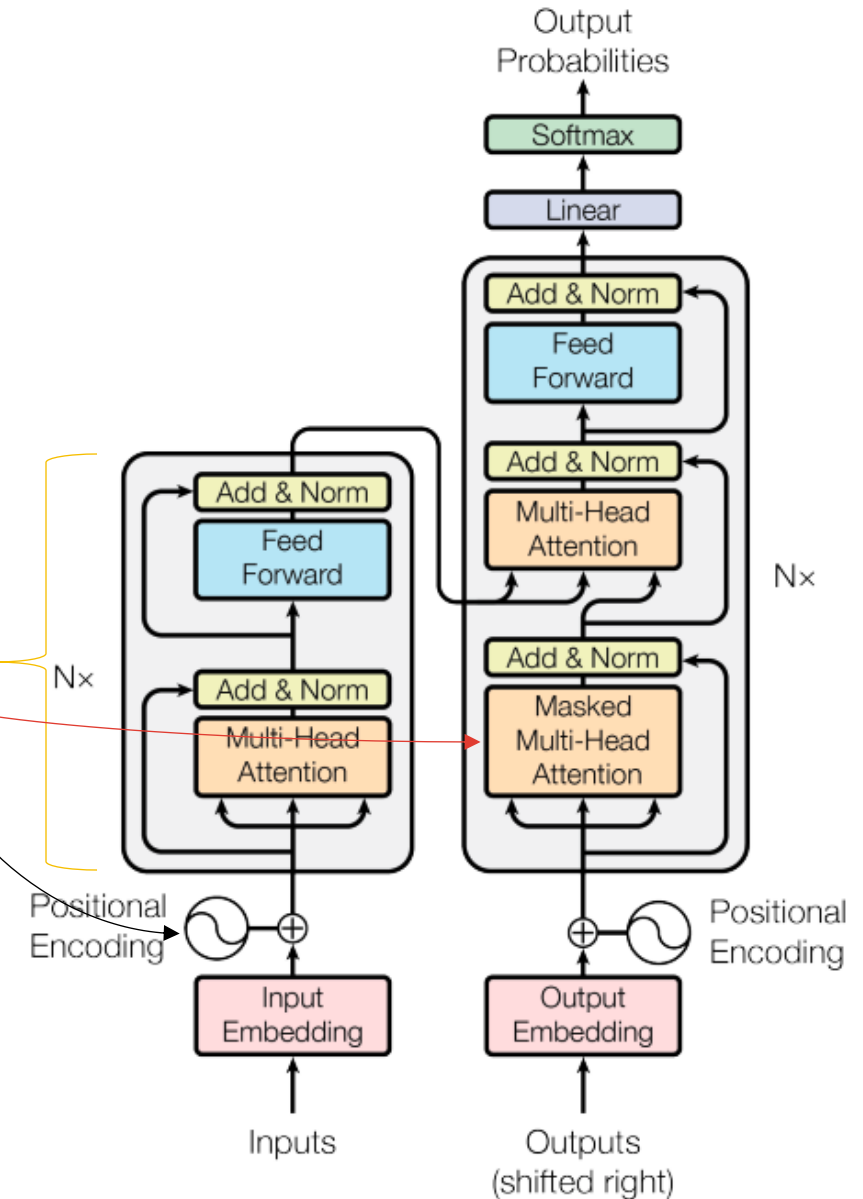
# Transformer

# Attention is all you need background

- Task: Translation
  - I.e., given a sentence in English, translate it to German
  - Training dataset contains **pairs** of (English, German)
  - At test time the model is provided an English and needs to **generate** a German translation
- Idea 
  - Jointly embed the input and (masked) output string in the same latent space
  - Use only **multi-head attention** and **(small) feed forward networks**

# Architecture

- Multi-head attention ✓
- Positional encoding/Input embeddings
- Encoder cell
- Masking in the decoder cell



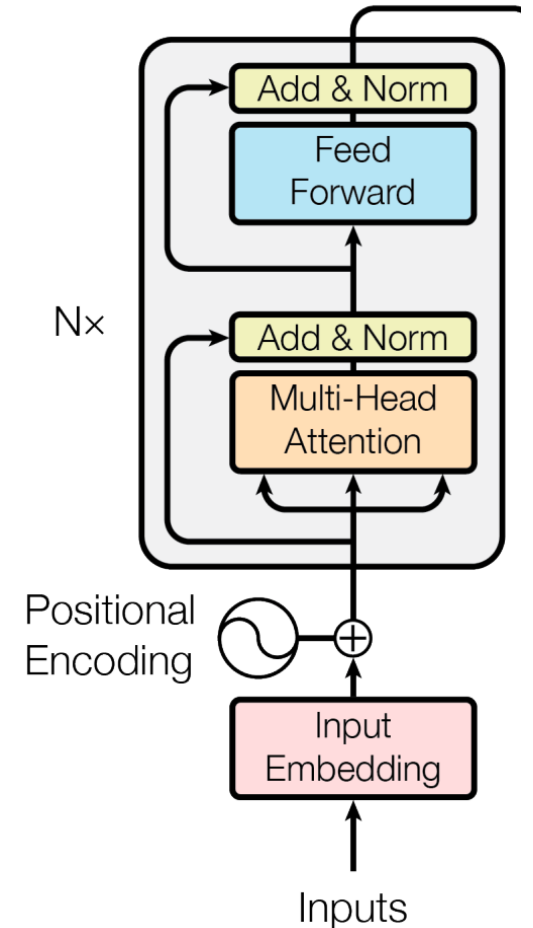
# Transformer inductive bias

- **Invariance**: the output of function  $f$  is **unaffected** by a “transformation” of the input
  - $f(\rho x) = f(x)$
- **Equivariance**: The input and output are **affected in the same way** by a transformation on the input
  - $f(\rho x) = \rho f(x)$
- Transformers define **positional invariance**
  - [“The”, “quick”] or [ “quick”, “The”] would define the **same** attention
  - This becomes less trivial with other contexts however, the principal remains



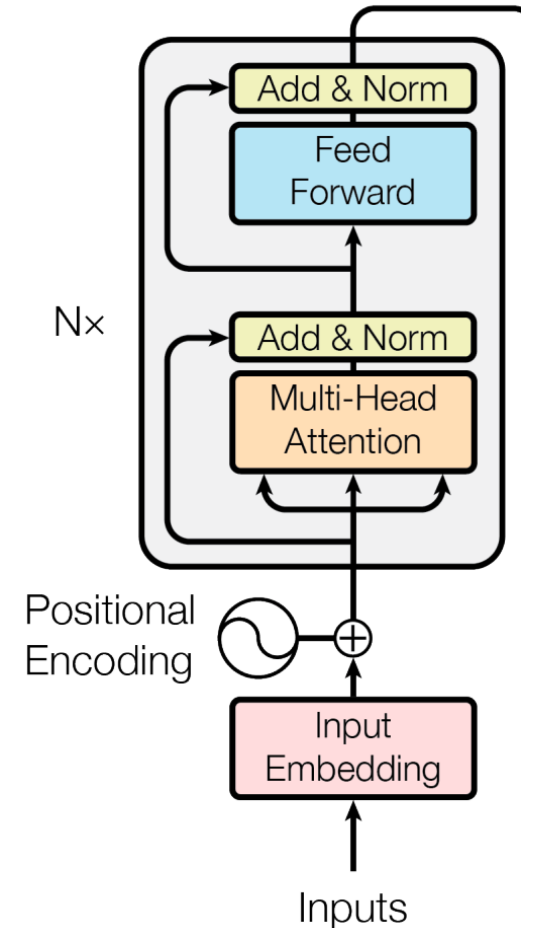
# Transformer positional invariance

- Transformers define **positional invariance**
  - ["The", "quick"] or [ "quick", "The"] would define the **same** attention
  - This becomes less trivial with other contexts however, the principal remains
- Positional invariance is **can be** useful however, for **sequences** we **don't want** positional invariance  $\Rightarrow$ 
  - Add **positional embeddings**
- Positional embeddings
  - $E_{(\text{pos}, 2i)} \sin\left(\frac{\text{pos}}{10000^{2i/d_X}}\right)$
  - $E_{(\text{pos}, 2i+1)} \cos\left(\frac{\text{pos}}{10000^{2i/d_X}}\right)$



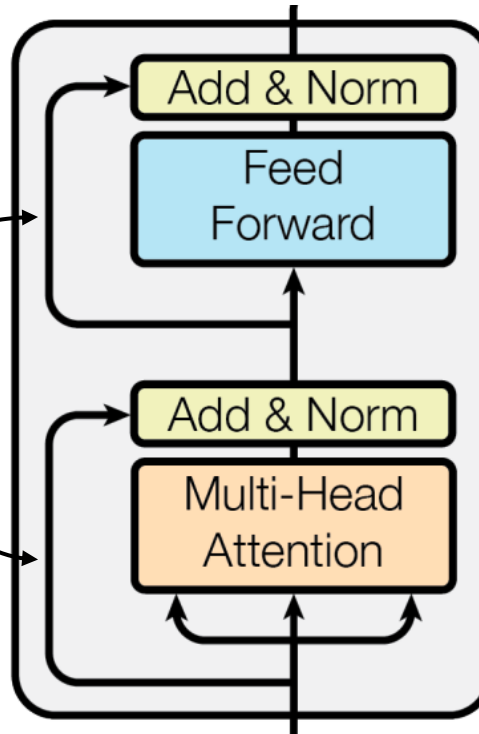
# Transformer positional invariance

- Positional embeddings
  - $E_{(\text{pos}, 2i)} \sin\left(\frac{\text{pos}}{10000^{2i/d_X}}\right)$
  - $E_{(\text{pos}, 2i+1)} \cos\left(\frac{\text{pos}}{10000^{2i/d_X}}\right)$
- The input to the multi-head attention becomes
  - Tokenize and convert to embeddings via lookup
  - $X = \begin{matrix} 0.5 & 0.2 \\ -1.2 & -1.4 \\ \vdots & \vdots \end{matrix}$  where  $X \in \mathbb{R}^{d_S \times d_X}$  and  $x$  is a row vector from  $X$
  - Obtain positional embeddings  $E \in \mathbb{R}^{1 \times d_X}$  (the positional embedding is constant across inputs)
  - $\tilde{X} = X + \tilde{E}$  where  $\tilde{E} \in \mathbb{R}^{d_S \times d_X}$  such that each **row vector** is equal to  $E$
- The Input embedding also contains a sentence identifier



# Transformer encoder cell

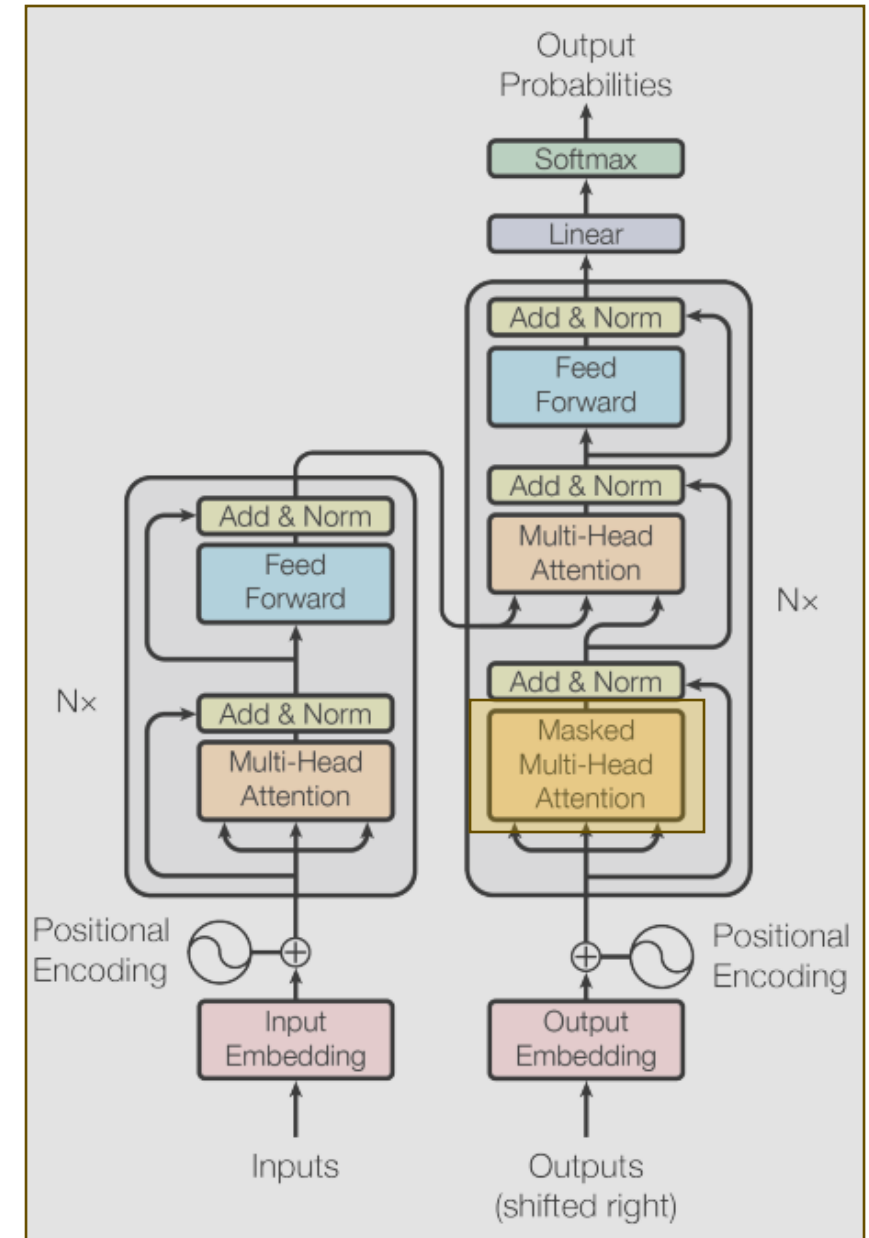
- Transformer cell contains:
  - Multi-head self attention
  - Generally small (2 layer) feed forward network, taking out =  $\text{concat}(\text{head}_1, \text{head}_2, \dots)W^O$  as an input
  - Residual connections



```
class EncoderLayer(nn.Module):  
  
    def __init__(  
        self,  
        d_model:int,  
        n_heads:int,  
        fc_dim:int,  
    ) -> None:  
        super().__init__()  
        self.fc_1 = nn.Linear(  
            d_model, fc_dim  
        )  
        self.relu = nn.ReLU()  
        self.fc_2 = nn.Linear(  
            fc_dim, d_model  
        )  
        self.mha = MultiHeadSelfAttention(  
            d_model=d_model, n_heads=n_heads  
        )  
        self.layer_norm_1 = nn.LayerNorm(d_model, eps=1e-6)  
        self.layer_norm_2 = nn.LayerNorm(d_model, eps=1e-6)  
  
    def forward(self, x:torch.Tensor):  
        mha_out = self.mha(x=x)  
        resid_norm = self.layer_norm_1(mha_out+x)  
        fc_out = self.fc_2(self.relu(self.fc_1(x)))  
        return self.layer_norm_2(fc_out+resid_norm)
```

# Masked attention

- Considering the test case:
  - Given an input in English: "I love ML"
  - Translate to German
  - **Importantly:** The model would not have access to the translated output
  - However, **as the model generates text**, it can bootstrap previous predictions i.e.
    - First generate "Ich"
    - To generate "liebe", the model can bootstrap its prediction of "Ich"

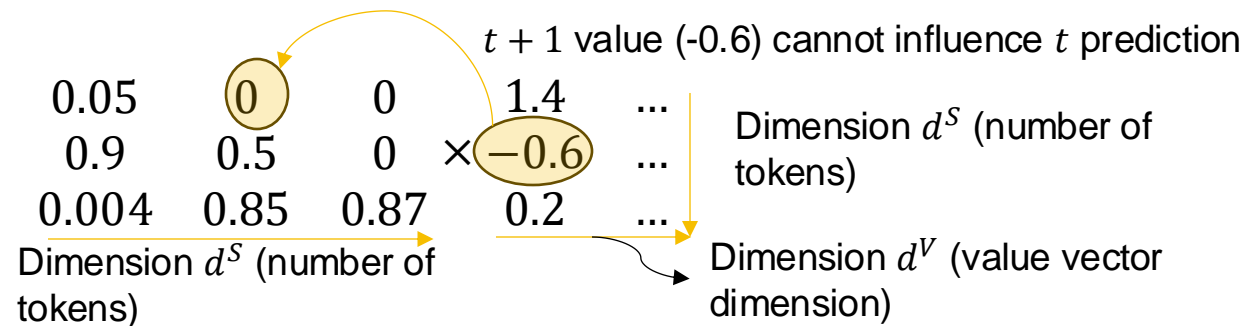


# Masked attention

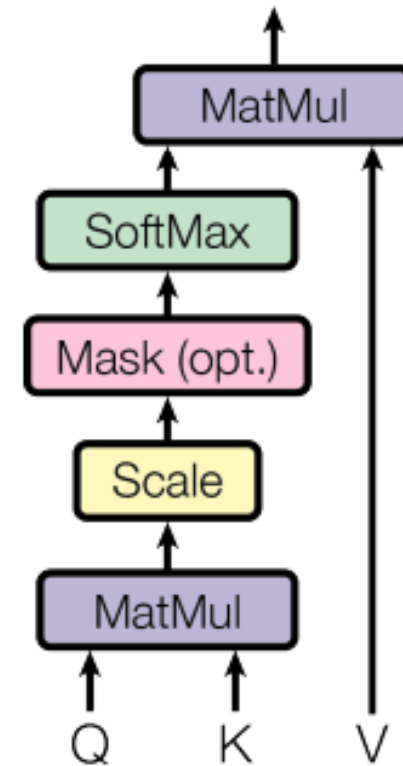
- Recall:

- $QK^T \in \mathbb{R}^{d_s \times d_s} = \begin{matrix} & \begin{matrix} 4.0 & -3.9 & 1.0 \end{matrix} \\ \begin{matrix} 1.8 \\ -0.6 \end{matrix} & \begin{matrix} -0.02 \\ 1.3 \end{matrix} & \begin{matrix} -2.0 \\ 1.4 \end{matrix} \end{matrix}$

- Mask the top triangle of the  $QK^T$  attention matrix, setting these values to  $-1e9$  (arbitrarily small value)
- When passed to the softmax, this will result in the following:



## Scaled Dot-Product Attention



# Training transformers

# Training transformers

- When training transformers:
  - Require **a lot** of data
  - **Larger batch sizes** (than usual) are required
  - Similarly, **lower learning rates** (than usual) are required
- Fundamentally: SGD/Adam is not **as well behaved** as other NN's (MLPs/RNNs/CNNs etc)
  - Require **tricks**

<https://arxiv.org/abs/1804.00247>

# LayerNorm

- Previously defined batch norm:

$$\text{BN}(x) = \frac{x - \mu(x)}{\sigma(x)} \beta + \gamma$$

- Where  $\mu$  and  $\sigma$  are defined **independently** for each **feature**
- Layer norm

$$\text{LN}(x) = \frac{x - \mu(x)}{\sigma(x)} \beta + \gamma$$

- Where  $\mu$  and  $\sigma$  are defined **independently** for each **sample**
- Why layer norm?
  - Does not require as large batch sizes (as required by BN)
  - More appropriate for batches with sequences, particularly of different lengths



# Attention is all you need tricks

- LayerNorm and Residual connections
- Adaptive learning rate:  $d_x^{-0.5} \min(s^{-0.5}, sw^{-1.5})$   
where  $s$  is algorithmic step and  $w$  is the number of warmup steps
  - Learning rate is proportional to the input embedding dimension  $d_x^{-0.5}$ . **Larger** the embedding, the **smaller** the learning rate
  - Learning rate increases linearly up to the warmup steps then decreases  $\propto s^{-0.5}$
- Regularisation:
  - Dropout
  - Label smoothing (adding  $\varepsilon$  probability mass to incorrect labels in classification)

# Attention is all you need tricks

- LayerNorm and Residual connections
  - ResNets enable **deeper** models
  - Normalisation stabilises the residual connections
  - <https://arxiv.org/pdf/1901.09321>
- Adaptive learning rate:  $d_x^{-0.5} \min(s^{-0.5}, sw^{-1.5})$  where  $s$  is algorithmic step and  $w$  is the number of warmup steps
  - Learning rate is proportional to the input embedding dimension  $d_x^{-0.5}$ . **Larger** the embedding, the **smaller** the learning rate
  - Learning rate increases linearly up to the warmup steps then decreases  $\propto s^{-0.5}$
- Regularisation:
  - Dropout
  - Label smoothing (adding  $\varepsilon$  probability mass to incorrect labels in classification)
- At the **start** of training transformers are **unstable**
  - Requiring **prohibitively** small learning rates
- Generally understood to be due to:
  - Layer Normalisation
  - (Potentially) amplified by Adam
- <https://proceedings.mlr.press/v119/huang20f.html>
- Prevent **overfitting**

# Understanding diffusion (through VAEs)

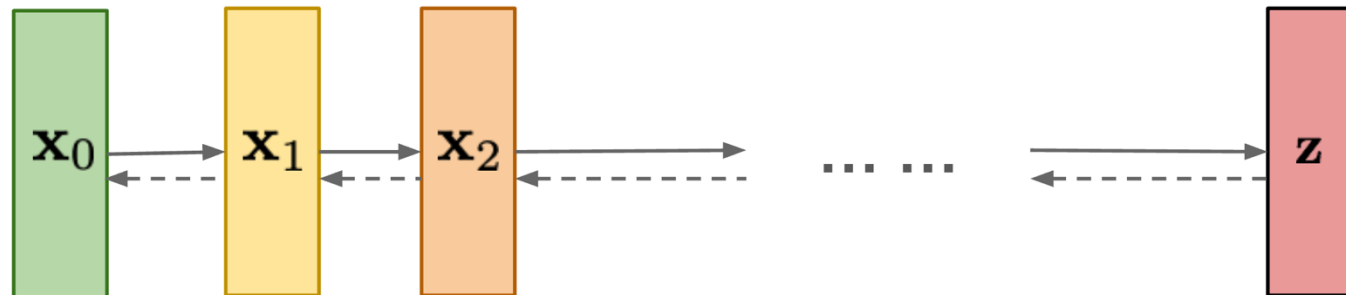
# Why diffusion models

- GANs previously held SOTA image generation however:
  - Are notoriously difficult to train
  - Have been demonstrated to cover less of the generation space in comparison to explicit likelihood models
- **Diffusion models overcome this** (<https://arxiv.org/pdf/2105.05233>)



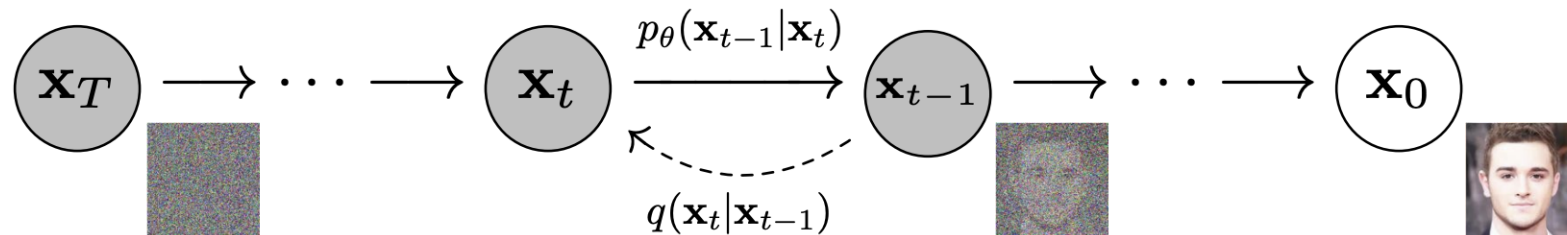
# Denoising diffusion probabilistic models intuition

- Gradually add gaussian noise and then learn a decoder to reverse noise into images
- Generate images by sampling from Gaussian noise



<https://arxiv.org/pdf/2006.11239>

# Denoising diffusion probabilistic models



- $p(x_0)$  defines the data distribution where  $x_0 \in \mathcal{X}_0$  and is of dimension  $d$
- $p_\theta(x_0) = \int p_\theta(x_{0:T}) dx_{1:T}$  is the proposed model where  $\mathcal{X}_{1:T}$  is of dimension  $d^T$  i.e., each  $x_i$  is a latent variable of the same dimension as  $x_0$ 
  - The same as VAEs,  $\theta$  are the model parameters and  $q$  is the **variational posterior**

# VAE loss vs Diffusion loss

$$\text{VAE, } p_{\theta}(x) = \int p_{\theta}(x|z)p_{\theta}(z)dz$$

$$\text{ELBO}(q_{\phi}(z|x)) = \mathbb{E}_{q_{\phi}(z|x)}(\log p_{\theta}(z, x)) - \mathbb{E}_{q_{\phi}(z|x)}(\log q_{\phi}(z|x)) \Rightarrow$$

$$L = \underbrace{-KL(q_{\phi}(z|x) || p_{\theta}(z))}_{\text{Posterior/Prior divergence}} + \underbrace{\mathbb{E}_{q_{\phi}(z|x)} [\log p_{\theta}(x|z)]}_{\text{Expected log likelihood}}$$

$$\text{Diffusion, } p_{\theta}(x) = \int p_{\theta}(x_0|x_{1:T})p_{\theta}(x_{1:T})dx_{1:T}$$

$$\text{ELBO}(q(x_{1:T}|x_0)) = \mathbb{E}_{q(x_{1:T}|x_0)}(\log p_{\theta}(x_{0:T})) - \mathbb{E}_{q(x_{1:T}|x_0)}(\log q(x_{1:T}|x_0)) \Rightarrow$$

Posterior/Prior divergence

Expected log likelihood

$$L = -KL(q(x_T|x_0) || p_{\theta}(x_T)) - \sum_{t>1} KL(q(x_{t-1}|x_t, x_0) || p_{\theta}(x_{t-1}|x_t)) + \mathbb{E}_{q(x_{1:T}|x_0)} [\log p_{\theta}(x_0|x_1)]$$

# Diffusion variational parameters

- $p_\theta(x) = \int p_\theta(x_0|x_{1:T})p_\theta(x_{1:T})dx_{1:T}$
- $L = -KL(q(x_T|x_0)||p_\theta(x_T)) - \sum_{t>1} KL(q(x_{t-1}|x_t, x_0)||p_\theta(x_{t-1}|x_t)) + \mathbb{E}_{q(x_{1:T}|x_0)} [\log p_\theta(x_0|x_1)]$

Where is  $\phi$ ?!





# Diffusion variational parameters

- $p_\theta(x) = \int p_\theta(x_0|x_{1:T})p_\theta(x_{1:T})dx_{1:T}$
- $L = -KL(q(x_T|x_0)||p_\theta(x_T)) - \sum_{t>1} KL(q(x_{t-1}|x_t, x_0)||p_\theta(x_{t-1}|x_t)) + \mathbb{E}_{q(x_{1:T}|x_0)} [\log p_\theta(x_0|x_1)]$
- “Forward process”/Variational posterior
  - $q(x_{1:T}|x_0) = \prod_{t=1}^T q(x_t|x_{t-1}) : q(x_t|x_{t-1}) \sim N(\sqrt{1 - \beta_t}x_{t-1}, \beta_t I)$
  - $\beta_1, \dots, \beta_T$  are the variational parameters **however**, these are considered **fixed** (in the simple case)
    - This is equivalent to a **very restrictive**  $\mathcal{F}$

# Diffusion loss decomposition

- $L = -KL(q(x_T|x_0)||p_\theta(x_T)) - \sum_{t>1} KL(q(x_{t-1}|x_t, x_0)||p_\theta(x_{t-1}|x_t)) + \mathbb{E}_{q(x_{1:T}|x_0)} [\log p_\theta(x_0|x_1)]$
- $p_\theta(x_{t-1}|x_t)$  is the **neural network we are training**
  - $p_\theta(x_{t-1}|x_t) = N(\mu_\theta(x_t, t), \Sigma_\theta(x_t, t))$  with  $\Sigma_\theta(x_t, t)$  fixed at  $\sigma_t^2 I$
  - Therefore, the neural network is modelling  $\mu_\theta(x_t, t)$
- $L_{t-1} \propto \frac{1}{2\sigma_t^2} \|\tilde{\mu}_t(x_t, x_0) - \mu_\theta(x_t, t)\|^2$ 
  - Where  $q(x_{t-1}|x_t, x_0) = N(\tilde{\mu}_t(x_t, x_0), \tilde{\beta}_t I)$  such that:
 

- $\tilde{\mu}_t(x_t, x_0) = \frac{\sqrt{\bar{\alpha}_{t-1}}\beta_t}{1-\bar{\alpha}_t} x_0 + \frac{\sqrt{\bar{\alpha}_t(1-\bar{\alpha}_t)}}{1-\bar{\alpha}_t} x_t$
    - $\tilde{\beta}_t = \frac{1-\bar{\alpha}_{t-1}}{1-\bar{\alpha}_t} \beta_t$

}

$\bar{\alpha}_t, \alpha_t$  and  $\tilde{\beta}_t$  are **reparameterisations** of the **fixed variational parameters**,  $\beta_t$

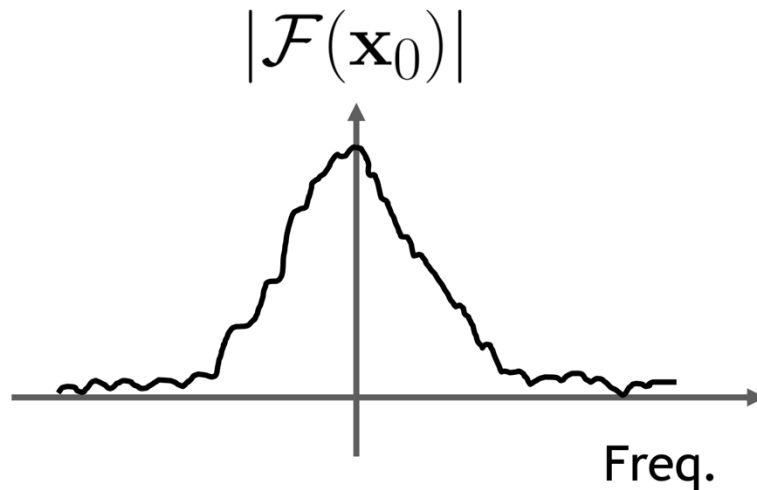
# Diffusion simple loss through reparameterisation

$$L = \mathbb{E}_{t, x_0, \epsilon} \left[ \underbrace{\frac{\beta_t^2}{2\sigma_t^2(1-\beta_t)(1-\bar{\alpha}_t)}}_{\lambda_t} \left\| \epsilon - \epsilon_\theta(\sqrt{\bar{\alpha}_t}x_0 + \sqrt{1-\bar{\alpha}_t}\epsilon, t) \right\|^2 \right]$$

- Where  $\epsilon$  is the noise added at time  $t$  and;
- $\epsilon_\theta$  is a **neural network** instead of  $\mu_\theta$
- The problem is **reparameterised** to predict the **noise** at step  $t$  rather than the **structure**
  - Empirically demonstrated to improve performance
- $\lambda_t$  is often **fixed at 1** regardless of the step
  - Produces a biased estimator **however** empirically this produces better results as it weights latent variables closer to  $x_t$  more greatly

# Understanding the forward process

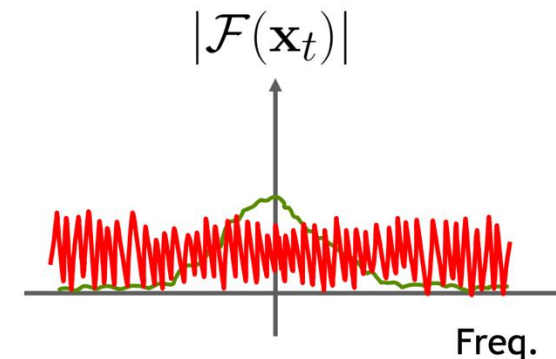
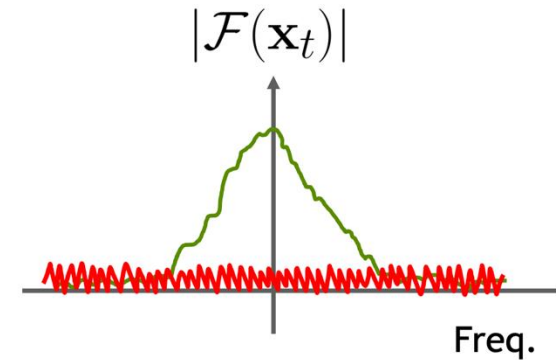
- Recall  $x_t = \sqrt{\bar{\alpha}_t}x_0 + \sqrt{1 - \bar{\alpha}_t}\varepsilon \Rightarrow$ 
  - $\mathcal{F}(x_t) = \sqrt{\bar{\alpha}_t}\mathcal{F}(x_0) + \sqrt{1 - \bar{\alpha}_t}\mathcal{F}(\varepsilon)$  where  $\mathcal{F}$  is the fourier transform



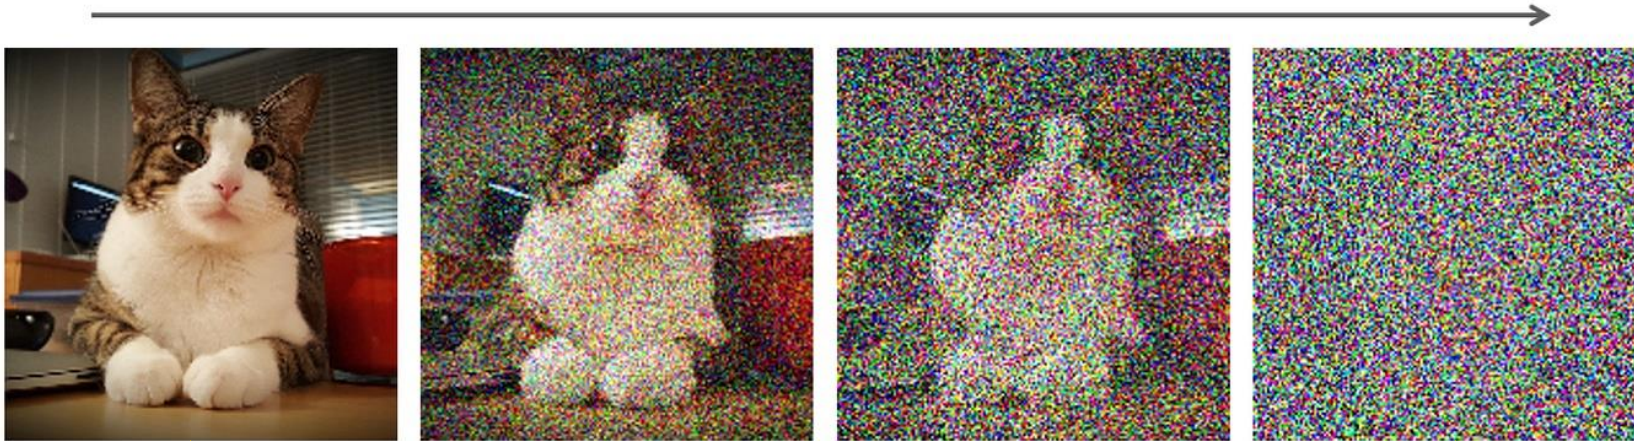
Small  $t$   
 $\bar{\alpha}_t \sim 1$



Large  $t$   
 $\bar{\alpha}_t \sim 0$



# Content-detail tradeoff



Denoising model is specialized for generating high-frequency content

Denoising model is specialized for generating low-frequency content

# Network architecture

- U-Net
  - Symmetric compression and decompression
  - Residual (as in ResNet) connections between compressed and decompressed layers of equal dimension
- Diffusion U-Net = U-Net +
  - Timestep embedding
  - Attention layer between during later compression/decompression stages

